

1. Creating the EKS service IAM role

- Create a EKSIAM role for VPC (option:Allows access to other AWS service resources that are required to operate clusters managed by EKS) and attach these 2 policies and administratoraccess

AmazonEKSClusterPolicy

AmazonEKSServicePolicy

- Create another EKSCluster role in the same way and keep only AmazonEKSClusterPolicy for Cluster Creation

2. Creating the VPC infrastructure using CloudFormation

- VPC, RouteTable, SecurityGroup, 3 Subnets

* Create stack with following s3 location as template, keep all the options as it is

<https://amazon-eks.s3.us-west-2.amazonaws.com/cloudformation/2020-10-29/amazon-eks-vpc-sample.yaml>

Note: While creating the stack, under the permissions page if the IAMRole created doesnt showup then

- Goto IAMRole, "Trust relationships" tab, If cloudformation.amazonaws.com isn't listed as a trusted entity in the bottom, then choose "Edit trust relationship"

- In the Policy Document editor, enter the following AWS CloudFormation service role trust policy:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "",
      "Effect": "Allow",
      "Principal": {
        "Service": "cloudformation.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

- Choose "Update Trust Policy"

Note: If you get error while creating VPC/IGW give full admin policy access to the IAM Role

3. Creating a ControlPlane & cluster in the AWS Management Console

- Choose the IAM Role / VPC / Subnet / SecuringGroup created above

4. Configuring kubectl for EKS

- Create a EC2 workstation for EKS

- Install Kubectl (<https://docs.aws.amazon.com/eks/latest/userguide/install-kubectl.html>)

```
$ curl -o kubectl https://amazon-eks.s3.us-west-2.amazonaws.com/1.17.11/2020-09-18/bin/linux/amd64/kubectl
```

```
$ chmod +x kubectl && sudo mv kubectl /usr/bin/kubectl
```

```
$ kubectl version --short --client
```

- Install AWS CLI

```
$ curl "https://awscli.amazonaws.com/awscli-exe-linux-x86_64.zip" -o "awscliv2.zip"
```

```
$ sudo apt install unzip && unzip awscliv2.zip
```

```
$ sudo ./aws/install --bin-dir /usr/bin --install-dir /usr/bin/aws-cli --update
```

```
$ aws --version
```

```
--- RUN as UBUNTU user -----
```

```
$ aws configure
```

```
$ aws eks list-clusters
```

```
$ aws eks update-kubeconfig --name wezvatechQA
```

Added new context arn:aws:eks:ap-south-1:184583194367:cluster/DemoEKSCluster to /home/ubuntu/.kube/config

```
$ kubectl config view
```

```
$ kubectl cluster-info
```

Kubernetes master is running at <https://4748266046766D1BC244740EE296AA39.sk1.ap-south-1.eks.amazonaws.com>

CoreDNS is running at <https://4748266046766D1BC244740EE296AA39.sk1.ap-south-1.eks.amazonaws.com/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy>

5. Provisioning Worker Nodes

- Choose the VPC, EKS Cluster, SG & 3 subnets created

<https://amazon-eks.s3.us-west-2.amazonaws.com/cloudformation/2020-10-29/amazon-eks-nodegroup.yaml>

- For NodeImageId, choose the AMI based on K8s version & AWS region (ami-0cd8562db082e8c1a)

<https://docs.aws.amazon.com/eks/latest/userguide/eks-optimized-ami.html>

(/aws/service/eks/optimized-ami/1.21/amazon-linux-2/recommended/image_id)

* Once the Nodes stack is created make a note of NodeInstanceRole from outputs: `arn:aws:iam::184583194367:role/DemoEKSNodes-NodeInstanceRole-1XFAAKMTKNJJ`

* Apply the aws-auth ConfigMap to your cluster (<https://docs.aws.amazon.com/eks/latest/userguide/add-user-role.html>)

```
$ curl -o aws-auth-cm.yaml https://amazon-eks.s3.us-west-2.amazonaws.com/cloudformation/2020-08-12/aws-auth-cm.yaml
```

```
curl -o aws-auth-cm.yaml https://amazon-eks.s3.us-west-2.amazonaws.com/cloudformation/2020-10-29/aws-auth-cm.yaml
```

- Edit the file & add the NodesInstanceRole ARN in it

```
$ kubectl apply -f aws-auth-cm.yaml
```

```
$ kubectl get nodes --watch
```

* Install Helm

```
$ curl https://get.helm.sh/helm-v3.2.3-linux-amd64.tar.gz > helm.tar.gz
```

```
$ tar xzvf helm.tar.gz
```

```
$ sudo mv linux-amd64/helm /usr/local/bin
```

===== [INGRESS Setup] =====

1. Create a IAM Policy

```
$ curl -o iam_policy.json https://raw.githubusercontent.com/kubernetes-sigs/aws-load-balancer-controller/v2.2.0/docs/install/iam_policy.json
```

```
$ aws iam create-policy \
  --policy-name AWSLoadBalancerControllerIAMPolicy \
  --policy-document file://iam_policy.json
```

2. Create a rbac-role.yml file with the following content & update with the above policy ARN at the end in Serviceaccount manifest

```
apiVersion: rbac.authorization.k8s.io/v1
```

```
kind: ClusterRole
```

```
metadata:
```

```
  labels:
```

```
    app.kubernetes.io/name: aws-load-balancer-controller
```

```
  name: aws-load-balancer-controller
```

```
rules:
```

```
- apiGroups:
```

```
  - ""
```

```
  - extensions
```

```
resources:
```

```

- configmaps
- endpoints
- events
- ingresses
- ingresses/status
- services
- pods/status
verbs:
- create
- get
- list
- update
- watch
- patch
- apiGroups:
  - ""
  - extensions
resources:
- nodes
- pods
- secrets
- services
- namespaces
verbs:
- get
- list
- watch
---
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  labels:
    app.kubernetes.io/name: aws-load-balancer-controller
  name: aws-load-balancer-controller
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: aws-load-balancer-controller
subjects:
- kind: ServiceAccount
  name: aws-load-balancer-controller
  namespace: kube-system
---
apiVersion: v1
kind: ServiceAccount
metadata:
  labels:
    app.kubernetes.io/name: aws-load-balancer-controller
  name: aws-load-balancer-controller
  namespace: kube-system
annotations:
  eks.amazonaws.com/role-arn: <your-iam-role-arn-for-alb-ingress-here>

$ kubectl apply -f rbac-role.yml

```

```
* Install eksctl on Ubuntu
$ curl --silent --location "https://github.com/weaveworks/eksctl/releases/latest/download/eksctl_$(uname -s)_amd64.tar.gz" | tar xz -C /tmp
$ sudo mv /tmp/eksctl /usr/local/bin
$ eksctl version
```

3. Create IAM OIDC provider:

```
$ eksctl utils associate-iam-oidc-provider \
  --region ap-south-1 \
  --cluster wezvatechQA \
  --approve
```

4. Create IAM Service Account & refer to the policy created at 1

```
$ eksctl create iamserviceaccount \
  --cluster wezvatechQA \
  --namespace kube-system \
  --name aws-load-balancer-controller \
  --attach-policy-arn arn:aws:iam::249397932281:policy/AWSLoadBalancerControllerIAMPolicy \
  --override-existing-serviceaccounts \
  --approve
$ eksctl get iamserviceaccount --cluster wezvatechQA
$ kubectl describe sa aws-load-balancer-controller -n kube-system
```

5. Install the TargetGroupBinding CRDs

```
$ kubectl apply -k github.com/aws/eks-charts/stable/aws-load-balancer-controller/crds?ref=master
$ kubectl get crd
```

6. Deploy AWS Load Balancer Controller (ALB Ingress Controller)

```
$ helm repo add eks https://aws.github.io/eks-charts
$ helm upgrade -i aws-load-balancer-controller \
  eks/aws-load-balancer-controller \
  -n kube-system \
  --set clusterName=wezvatechQA \
  --set serviceAccount.create=false \
  --set serviceAccount.name=aws-load-balancer-controller \
```

```
$ kubectl -n kube-system rollout status deployment aws-load-balancer-controller
$ kubectl get pods -n kube-system
```

7. Deploy sample application

```
$ curl -o demoapp.yml https://raw.githubusercontent.com/kubernetes-sigs/aws-load-balancer-controller/v2.2.0/docs/examples/2048/2048_full.yaml
-- change the replica to 2
$ kubectl apply -f demoapp.yml
$ kubectl get ingress/ingress-2048 -n game-2048
```

---CLEANUP-----

```
$ kubectl delete -f demoapp.yml
$ helm uninstall aws-load-balancer-controller -n kube-system
$ kubectl delete -k github.com/aws/eks-charts/stable/aws-load-balancer-controller/crds?ref=master
$ eksctl delete iamserviceaccount \
  --cluster wezvatechQA \
  --name aws-load-balancer-controller \
  --namespace kube-system \
  --wait
```

```
$ aws iam delete-policy \
--policy-arn arn:aws:iam::249397932281:policy/AWSLoadBalancerControllerIAMPolicy
```

===== [Cluster Backup] =====

1. Create a S3 bucket

2. Install Velero Client

```
$ wget https://github.com/vmware-tanzu/velero/releases/download/v1.3.2/velero-v1.3.2-linux-amd64.tar.gz
```

```
$ tar -xvf velero-v1.3.2-linux-amd64.tar.gz -C /tmp
```

```
$ sudo mv /tmp/velero-v1.3.2-linux-amd64/velero /usr/local/bin
```

```
$ velero version
```

3. Install Velero on EKS

```
$ velero install \
```

```
--provider aws \
```

```
--plugins velero/velero-plugin-for-aws:v1.0.1 \
```

```
--bucket <S3BucketName> \
```

```
--backup-location-config region=ap-south-1 \
```

```
--snapshot-location-config region=ap-south-1 \
```

```
--secret-file /home/ubuntu/.aws/credentials
```

```
$ kubectl logs -f deployment/velero -n velero
```

4. Deploy test app

```
$ kubectl create ns demo
```

```
$ kubectl apply -f testpod.yml
```

5. Create Backup

```
$ velero backup create backup.1 --include-namespaces demo
```

```
$ velero backup describe backup.1
```

- Verify the backup folder inside s3 bucket

```
$ velero backup get
```

6. Create a disaster scenario i.e delete namespace

```
$ kubectl delete ns demo
```

7. Restore from backup

```
$ velero restore create --from-backup backup.1
```

```
$ kubectl get ns
```

----CleanUp----

```
$ velero backup delete backup.1
```

```
$ kubectl delete ns demo
```

===== [CLUSTER AUTOSCALER SETUP] =====

1. Deploy Metrics server

```
$ kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/download/v0.5.0/components.yaml
```

```
$ kubectl get pods -n kube-system
```

```
$ kubectl top pods or kubectl top nodes
```

2. Configure Autoscaling Group

- Make a note of existing Autoscaling Group created by EKS Worker Stack

- Edit IAM role associated to EKS node & add an inline IAM policy with the following document

```
{
```

```

"Version": "2012-10-17",
"Statement": [
  {
    "Action": [
      "autoscaling:DescribeAutoScalingGroups",
      "autoscaling:DescribeAutoScalingInstances",
      "autoscaling:DescribeLaunchConfigurations",
      "autoscaling:DescribeTags",
      "autoscaling:SetDesiredCapacity",
      "autoscaling:TerminateInstanceInAutoScalingGroup",
      "ec2:DescribeLaunchTemplateVersions"
    ],
    "Resource": "*",
    "Effect": "Allow"
  }
]
}

```

3. Download the yaml & add the Cluster name at the end, apply the yaml

```
$ wget https://raw.githubusercontent.com/kubernetes/autoscaler/master/cluster-autoscaler/cloudprovider/aws/examples/cluster-autoscaler-autodiscover.yaml
```

```
- --nodes=1:6:<Autoscaling Group Name>
```

```
$ kubectl apply -f cluster-autoscaler-autodiscover.yaml
```

```
$ kubectl get pods -n kube-system
```

```
$ kubectl logs -f <podname> -n kube-system
```

4. Deploy sample application and increase the replicas to add load

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: nginx-to-scaleout
```

```
spec:
```

```
  replicas: 1
```

```
  selector:
```

```
    matchLabels:
```

```
      app: nginx
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        service: nginx
```

```
        app: nginx
```

```
    spec:
```

```
      containers:
```

```
      - image: nginx
```

```
        name: nginx-to-scaleout
```

```
      resources:
```

```
        limits:
```

```
          cpu: 500m
```

```
          memory: 512Mi
```

```
        requests:
```

```
          cpu: 500m
```

```
          memory: 512Mi
```

---CLEANUP----

```
$ kubectl delete -f cluster-autoscaler-autodiscover.yaml
```

- Remove the inline policy from the IAM Role for the worker node group

\$ kubectl delete -f <https://github.com/kubernetes-sigs/metrics-server/releases/download/v0.5.0/components.yaml>